Rowing Speed Prediction from Per-Stroke Telemetry Data

Nate Ly

May 4, 2025

Introduction

In competitive rowing, the goal is to travel a distance of 2000 meters as fast as possible. The fastest boat, called an Eight, has eight rowers and a coxswain. Boat speed comes from a combination of power output from the athletes and technical cohesion. While maximizing power output relies on training the human body to its physiological limits, technical cohesion depends on neuromuscular synchronization and practice. Although sports science has made great strides in optimizing human physiology through training and nutrition, rowing technique remains the bottleneck of boat speed.

Traditional fluid-dynamics models can predict boat motion, but they require extreme care in specifying boundary conditions and are computationally expensive. Moreover, they cannot capture human variability due to fatigue or subtle technical differences. A data-driven approach "priced in" these unmodeled factors, offering not only a cheaper but potentially more practical and insightful method.

General Approach:

- Collect per-stroke data through special oarlocks equipped with strain gauges.

- Clean and filter the data to retain only "quality" strokes—i.e., when all eight rowers are rowing with intent and the boat is neither spinning nor drilling.

- Train a neural network to map the extracted features to boat speed.

Challenges:

- The telemetry oarlocks are in beta testing, endure substantial and repetitive strain, and are calibrated by hand, which can introduce skew in the measurements.

- Determining the best method to filter for "quality" strokes.

- Tuning machine-learning hyperparameters to optimize loss and generalization.

Data

The data were collected from Yale Heavyweight Crew on April 7, 2025, using telemetry from two Eights, with the written permission of all 16 rowers and the head coach. Each per-stroke record includes:

- Power (W)

- Rate (strokes per minute)

- Work (J)

- Blade angles: catch angle, slip, wash, finish angle, connected length (degrees)

- Peak force angle (degrees)

- Catch time, stroke duration (seconds)

The target variable, boat speed (m/s), was measured via GPS.

Methodology

Problem Formulation

Let each stroke be represented by a feature vector $\mathbf{x} \in \mathbb{R}^d$ and its measured speed by $y \in \mathbb{R}$. Introduce a fixed offset $b_0 \in \mathbb{R}$ to account for baseline speed. We then learn a residual function

$$g : \mathbb{R}^d \rightarrow \mathbb{R}$$

such that

$$g(\mathbf{x}) \approx y - b_0.$$

The full predictor is

$$\hat{y} = g(\mathbf{x}) + b_0.$$

Model Architecture

We implement a feed-forward neural network in PyTorch with four fully connected layers and ReLU activations. The layer dimensions are $d \to 64 \to 32 \to 16 \to 1$. Let $\mathbf{h}^{(0)} = \mathbf{x}$. Then for $i = 1, 2, 3$:

$$\mathbf{z}^{(i)} = w^{(i)}\,\mathbf{h}^{(i-1)} + \mathbf{b}^{(i)}$$

$$\mathbf{h}^{(i)} = \text{ReLU}(\mathbf{z}^{(i)})$$

$$\mathbf{z}^{(4)} = w^{(4)}\,\mathbf{h}^{(3)} + b^{(4)}, \quad \hat{y} = \mathbf{z}^{(4)}.$$

Loss Function and Optimization

We aim to minimize the Mean Squared Error over n strokes:

$$\mathcal{L}(\theta) = \frac{1}{N} \sum_{i=1}^{N} \big(f(\mathbf{x}_i) - (y_i)\big)^2.$$

Gradients are computed via back-propagation and weights are adjusted using PyTorch's Adam (Adaptive Moment Estimation)

- $g^{(t)} = \nabla_w L^{(t)} = \frac{\partial \text{Loss}}{\partial \text{Weights}}$: Gradient of the loss with respect to weights at time step t

- $m^{(t)}$: First momentum (expected value) with decay β_1

- $v^{(t)}$: Second momentum (variance) with decay β_2

$$m^{(t)} = \beta_1 \cdot m^{(t-1)} + (1 - \beta_1) \cdot g^{(t)}$$

$$v^{(t)} = \beta_2 \cdot v^{(t-1)} + (1 - \beta_2) \cdot \big(g^{(t)}\big)^2$$

$$m^{(t)} = \frac{m^{(t)}}{1 - \beta_1^t}, \quad v^{(t)} = \frac{v^{(t)}}{1 - \beta_2^t}$$

For learning rate α and small constant ϵ to avoid division by zero upon convergence, the weights are updated as follows:

$$w^{(t+1)} = w^{(t)} - \alpha \cdot \frac{m^{(t)}}{\sqrt{v^{(t)}} + \epsilon}$$

Implementation

Data Processing

To retain only quality strokes, any strokes below a certain speed were filtered out. Given the features listed before, the additional variance feature between the eight rowers of each feature on the crew was added. Additionally, power and work were cube root transformed as they are cubically related.

Parameters

- **Minimum Speed(s)** 4.75, 5

- **Epochs** 2000

- **Weight Decay** 1×10^{-7}

- **Learning Rate** 5×10^{-4}

- **Batch Size** 16

- **Momentum** N/A (using default Adam with $\beta_1 = 0.9$, $\beta_2 = 0.999$)

- **Dropout Rate** 0.2

- **Training Split** 80% train / 20% test

Results

The original goal was to obtain a model that would predict speed of "at-pace", or hard rowing strokes to 1s/500m, around a 1 percent mean absolute percent error. Below is a reference chart for different rowing speeds and relative efforts

Table 1: Rowing Speeds and Corresponding 500 m Paces and Intensities

Speed (m/s)	Pace per 500 m	Intensity
4.50	1:51.1	Low
4.75	1:45.0	Steady State
5.00	1:40.0	Medium rowing
5.25	1:35.0	Hard rowing
5.50	1:31.0	Around Race Effort

Minimum Speed of 4.75 m/s

Filtering out any non-"quality" strokes (slower than a 1:45 pace), we obtain the following results:

Table 2: Model performance metrics for strokes at or above 4.75 m/s

| Metric	Value
Number of data points	1678
MSE	0.0048
R^2 score	0.8261
MAPE	1.02%
Split error over 500 m	1.00 s

Figure 1: MSE Loss vs. Epochs

Figure 2: Scatterplot for strokes above $4.75\,\mathrm{m/s}$

Minimum Speed of 5 m/s

If we restrict strokes to moderate effort or more, we can further reduce noise while still using data that encompasses the "hard" effort strokes.

Table 3: Model performance metrics for the filtered dataset

| Metric	Value
Number of data points	1057
MSE	0.0017
R^2 score	0.7511
MAPE	0.68%
Split error over 500 m	0.65 s

Figure 3: Scatterplot for strokes above 5 m/s

Extra: Minimum Speed of 5 m/s with Previous Speed Data

If we add the prior stroke's speed to our inputs, the features represent speed delta rather than absolute speed

Table 4: Model performance metrics for the filtered dataset

| Metric	Value
Number of data points	1055
MSE	0.0013
R^2 score	0.8056
MAPE	0.53%
Split error over 500 m	0.56 s

Figure 4: Scatterplot for strokes above 5 m/s with previous stroke speed

Conclusion

In conclusion, we were successful in our target of predicting boat speed to under a 1 split error over 500m. The data was filtered, and features were engineered that led to a successful neural network. Further improvements could include using a grid search for parameter tuning, and collecting more quality data to reduce the risk of overfitting.